




p2p architecture  
peer to peer.

Client-Server:

mail.

chat appl (peers)



P2P.

→ absence of infrastructure  
(Server)

→ "distributed" → truly

Server

always on  
sufficient

bandwidth

Comp. resources


P2P

Chat Appl.

→ no server. → owner ✓  
→  
diff. for "authorities"  
to control apps.

---


p2p.



$N = \text{nr}$   
of nodes

- Manually set up N/w
  - Small # of nodes.
- N

Cl<sub>1</sub> - Ser.



apls  
phy n/w.

logical links

# P2P

1. How to form p2p network:-

2. Resources are also distributed

3. Distributed Search Algo

→ Brute force approach →

$n \rightarrow$  all others


expensive.

$$n^2$$

→ Scheme for distribution:-

$10^6$  users

→ request  
proc. ~~other~~  
b/w.




# Pure P2P architecture

- *no* always-on server
- arbitrary end systems directly communicate
- peers are intermittently connected and change IP addresses

## *examples:*

- file distribution (BitTorrent)
- Streaming (KanKan)
- VoIP (Skype)

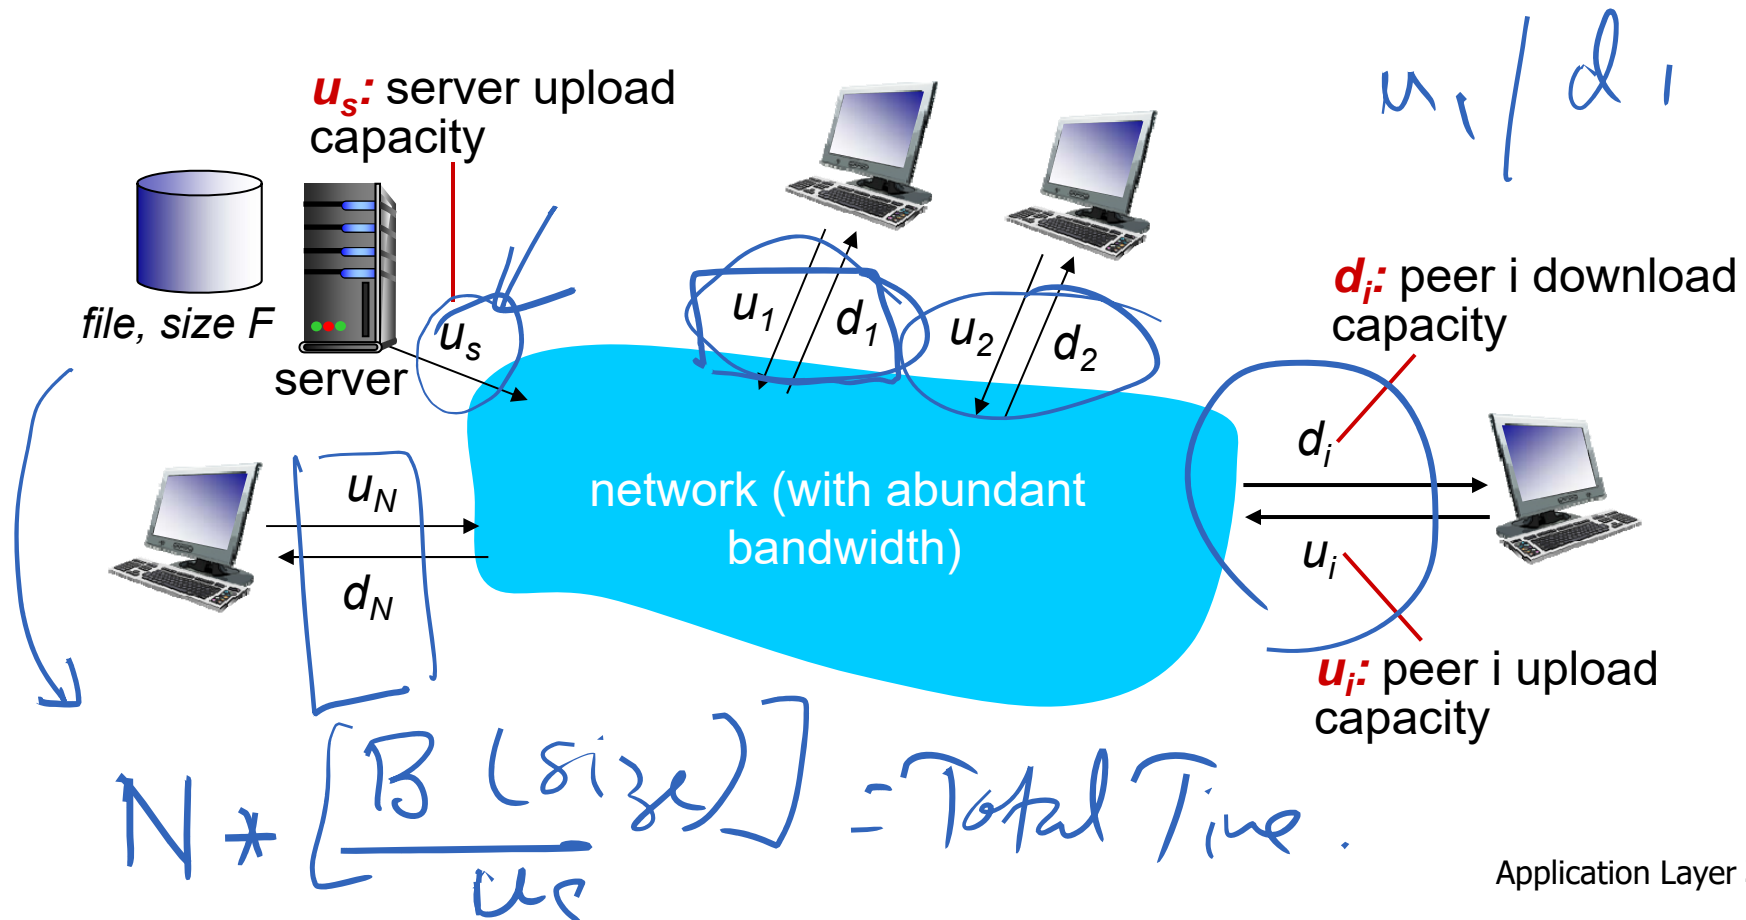




# File distribution: client-server vs P2P


Question: how much time to distribute file (size  $F$ ) from one server to  $N$  peers?

- peer upload/download capacity is limited resource



# File distribution time: client-server

- **server transmission:** must sequentially send (upload)  $N$  file copies:
  - time to send one copy:  $F/u_s$
  - time to send  $N$  copies:  $NF/u_s$
- **client:** each client must download file copy
  - $d_{min}$  = min client download rate
  - min client download time:  $F/d_{min}$




*time to distribute  $F$   
to  $N$  clients using  
client-server approach*

$$D_{c-s} \geq \max\{NF/u_s, F/d_{min}\}$$

increases linearly in  $N$

# File distribution time: P2P

- **server transmission:** must upload at least one copy
  - time to send one copy:  $F/u_s$
- **client:** each client must download file copy
  - min client download time:  $F/d_{\min}$
- **clients:** as aggregate must download  $NF$  bits
  - max upload rate (limiting max download rate) is  $u_s + \sum u_i$



time to distribute  $F$   
to  $N$  clients using  
P2P approach


$$D_{P2P} \geq \max\{F/u_s, F/d_{\min}, NF/(u_s + \sum u_i)\}$$

increases linearly in  $N$  ...

... but so does this, as each peer brings service capacity

# Client-server vs. P2P: example

client upload rate =  $u$ ,  $F/u = 1$  hour,  $u_s = 10u$ ,  $d_{min} \geq u_s$



# P2P file distribution: BitTorrent


Cooperation

- file divided into 256Kb chunks
- peers in torrent send/receive file chunks

**tracker:** tracks peers participating in torrent


**torrent:** group of peers exchanging chunks of a file

Alice arrives ...  
... obtains list  
of peers from tracker  
... and begins exchanging  
file chunks with peers in torrent



# P2P file distribution: BitTorrent

- peer joining torrent:
  - has no chunks, but will accumulate them over time from other peers
  - registers with tracker to get list of peers, connects to subset of peers (“neighbors”)
- while downloading, peer uploads chunks to other peers
- peer may change peers with whom it exchanges chunks
- **churn**: peers may come and go
- once peer has entire file, it may (selfishly) leave or (altruistically) remain in torrent



# Network Function.


protocol.

P2P → "Chat"

↙  
Cafeteria → sitting

→ underlying protocol.

Assignment




P2P


- Formation of P2P N/w
- Resource distribution and Search
- Cooperation

---

P2P N/w (Overlay n/w)







# Hash Table

Name  
 $f(\text{Name}) \rightarrow$  1st ch  
26 ch:

$$f(\underset{\text{key}}{x}) = h.$$

Search space

Key Space

$\& (x)$



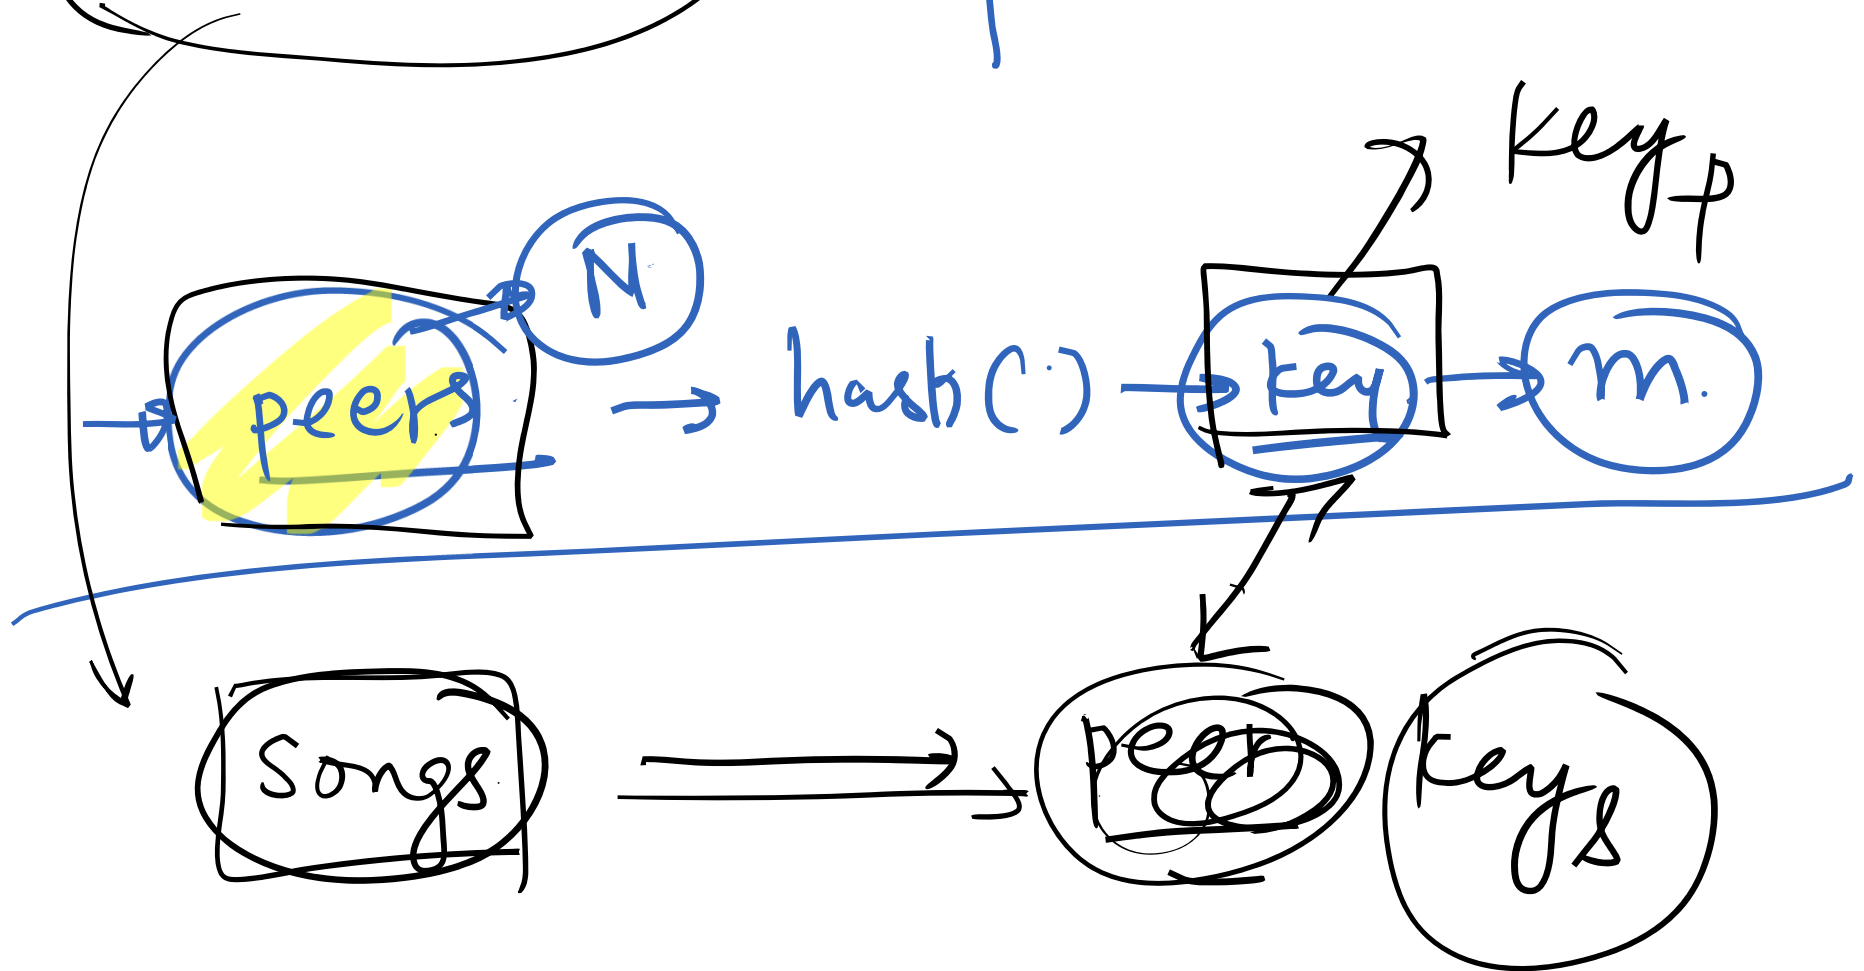
1-emb 2's

Sanjny

$\rightarrow \text{key}[S]$

Distribution

Search.



Cooperation


file Transfer

Token based  
Crp.

TIT for TAT strategy.

Generates TFT.

1st time  
always Crp.



# BitTorrent: requesting, sending file chunks

## *requesting chunks:*


- at any given time, different peers have different subsets of file chunks
- periodically, Alice asks each peer for list of chunks that they have
- Alice requests missing chunks from peers, rarest first

## *sending chunks: tit-for-tat*

- Alice sends chunks to those four peers currently sending her chunks *at highest rate*
  - other peers are choked by Alice (do not receive chunks from her)
  - re-evaluate top 4 every 10 secs
- every 30 secs: randomly select another peer, starts sending chunks
  - “optimistically unchoke” this peer
  - newly chosen peer may join top 4

# BitTorrent: tit-for-tat

- (1) Alice “optimistically unchokes” Bob
- (2) Alice becomes one of Bob’s top-four providers; Bob reciprocates
- (3) Bob becomes one of Alice’s top-four providers




P2P

architecture

Scaling

~~server specific~~

distributed load.








Anchor (node) → maintain a list of neighbors

join req. → accept  
 → send a random neighbor id.



pzp :

node : resource constraint  
power constraint

rand()

p2p N/w.


anchor.

join.

prop. to "all"

flooding

for peer  
p2p nodes.




relay

if ( ~~power~~ Energy left < threshold )  
don't ~~can~~ accept

if ( deg(node) > deg\_max )  
✓  
don't accept

Anchor  $\rightarrow$  nucleus.



How to ensure connected  
ness

2, 5, 6, 7

→ neighbour of neighbour  
(2-hop) list

---

N/w function


→ Anchor node (s) (well known)

→ Conn. ref → Arch.

p( ) rand( ) accept  
id of a neigh

→ ~~o~~ n/w partition → fix that

resource dist. & search



resource  $\rightarrow$  distribute it in p2p

song  $\rightarrow$   $\left[ \begin{array}{l} s_1 \rightarrow \text{node } j \\ s_2 \rightarrow \text{node } k \end{array} \right]$





# Chapter 2: outline

2.1 principles of network applications

2.2 Web and HTTP

2.3 electronic mail

- SMTP, POP3, IMAP

2.4 DNS

2.5 P2P applications

2.6 video streaming and content distribution networks (CDNs)

2.7 socket programming with UDP and TCP

# Multimedia Streaming.

Real Time Application

X — delay guarantees

MM  $\rightarrow$  bandwidth constraint

+ Client-Server  $\rightarrow$  requirements  
are very heavy

P2P

{ a cooperation problem / Monitored problem  
availability

problem. stay within  
Client-Server model and  
provide scalable streaming  
service.

---

Multiple Servers.

→ load balancing

→ storage of resources.

→ google.com → US/Korea/India  
maps to one of the multiple server

- DNS DB has list of mappings  
+ clients (asks for a map)

↳ client IP addr →


local

IP addr → locally mapping

→ Server

server  $\rightarrow$  replicated (multiple)

resources  $\rightarrow$  distributed



server [Table]

| movie1  | server 4 |
|---------|----------|
| movie 2 | server 5 |

Server.

Table.

| Movie key | Server IDs |
|-----------|------------|
| movies    | S1, S3, S9 |

Client  $\text{req}(\text{movie}_1) \rightarrow S1$

Server S1:  $\text{resp}(\text{movie}_1 \rightarrow \boxed{S1, S3, S9})$

Client  $\rightarrow \text{req}(\text{movie}_1, S3)$

# MM distribution Resolution:

Client

✓ device Cap  
mobile old  
480p

BW

Movie →

diff. versions

at

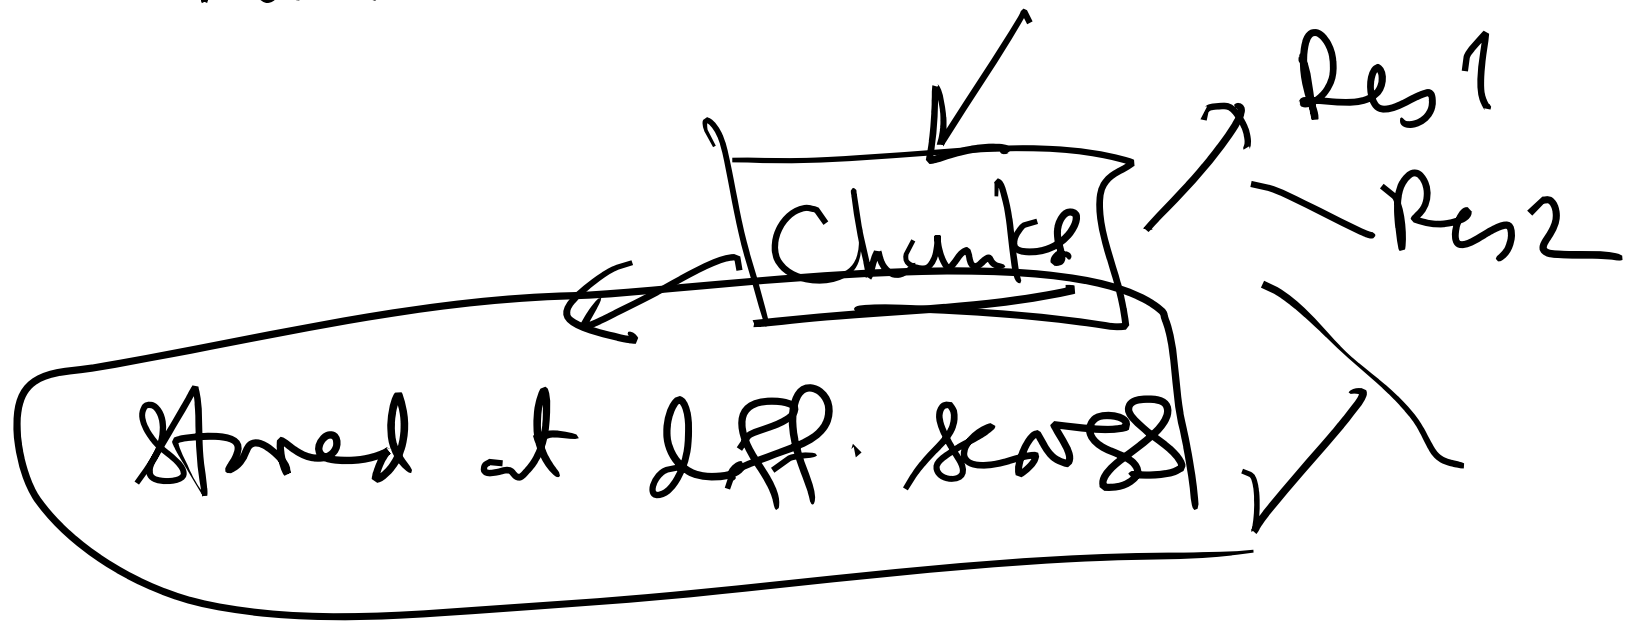
each resolution

~~Fixed~~

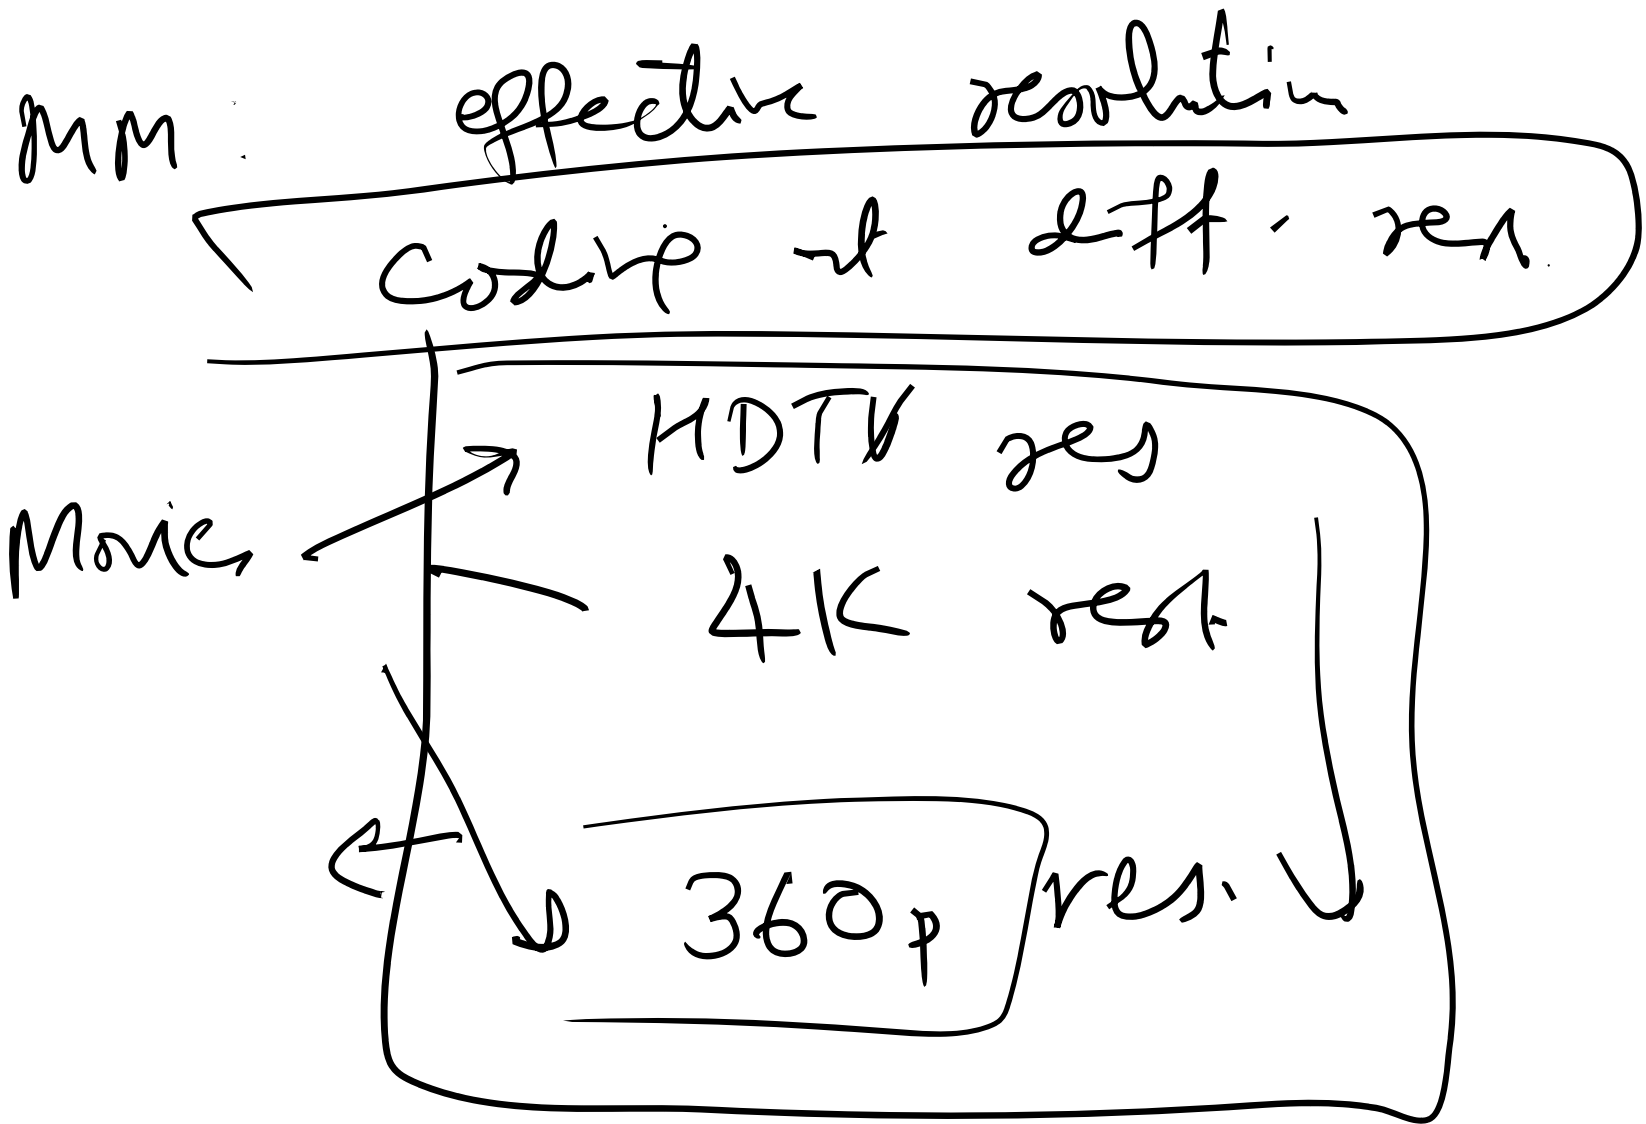
dynamic.

→ runtime

movie  $\rightarrow$  vesting at res.







# Video Streaming and CDNs: context


- video traffic: major consumer of Internet bandwidth
  - Netflix, YouTube: 37%, 16% of downstream residential ISP traffic
  - ~1B YouTube users, ~75M Netflix users
- challenge: scale - how to reach ~1B users?
  - single mega-video server won't work (why?)
- challenge: heterogeneity
  - different users have different capabilities (e.g., wired versus mobile; bandwidth rich versus bandwidth poor)
- **solution:** distributed, application-level infrastructure



# Multimedia: video

- video: sequence of images displayed at constant rate
  - e.g., 24 images/sec
- digital image: array of pixels
  - each pixel represented by bits
- coding: use redundancy *within* and *between* images to decrease # bits used to encode image
  - spatial (within image)
  - temporal (from one image to next)

*spatial coding example:* instead of sending  $N$  values of same color (all purple), send only two values: color value (*purple*) and number of repeated values ( $N$ )



frame  $i$

*temporal coding example:* instead of sending complete frame at  $i+1$ , send only differences from frame  $i$



frame  $i+1$




# Multimedia: video

- **CBR: (constant bit rate):**  
video encoding rate fixed
- **VBR: (variable bit rate):**  
video encoding rate changes  
as amount of spatial,  
temporal coding changes
- **examples:**

- MPEG I (CD-ROM) 1.5 Mbps
- MPEG2 (DVD) 3-6 Mbps
- MPEG4 (often used in Internet, < 1 Mbps)

avg Rate

*spatial coding example:* instead of sending  $N$  values of same color (all purple), send only two values: color value (purple) and number of repeated values ( $N$ )



frame  $i$




frame  $i+1$

*temporal coding example:* instead of sending complete frame at  $i+1$ , send only differences from frame  $i$

# Streaming stored video:

simple scenario:



# Streaming multimedia: DASH

- **DASH**: **D**ynamic, **A**daptive **S**treaming over **H**TTP

- **server**:

- divides video file into multiple chunks
- each chunk stored, encoded at different rates

- **manifest file**: provides URLs for different chunks

- **client**:

- periodically measures server-to-client bandwidth
- consulting manifest, requests one chunk at a time
  - chooses maximum coding rate sustainable given current bandwidth
  - can choose different coding rates at different points in time (depending on available bandwidth at time)

# Streaming multimedia: DASH

- *DASH: Dynamic, Adaptive Streaming over HTTP*
- “intelligence” at client: client determines
  - *when* to request chunk (so that buffer starvation, or overflow does not occur)
  - *what encoding rate* to request (higher quality when more bandwidth available)
  - *where* to request chunk (can request from URL server that is “close” to client or has high available bandwidth)

# Content distribution networks

- *challenge*: how to stream content (selected from millions of videos) to hundreds of thousands of *simultaneous* users?
- *option 1*: single, large “mega-server”
  - single point of failure
  - point of network congestion
  - long path to distant clients
  - multiple copies of video sent over outgoing link

....quite simply: this solution *doesn't scale*




# Content distribution networks

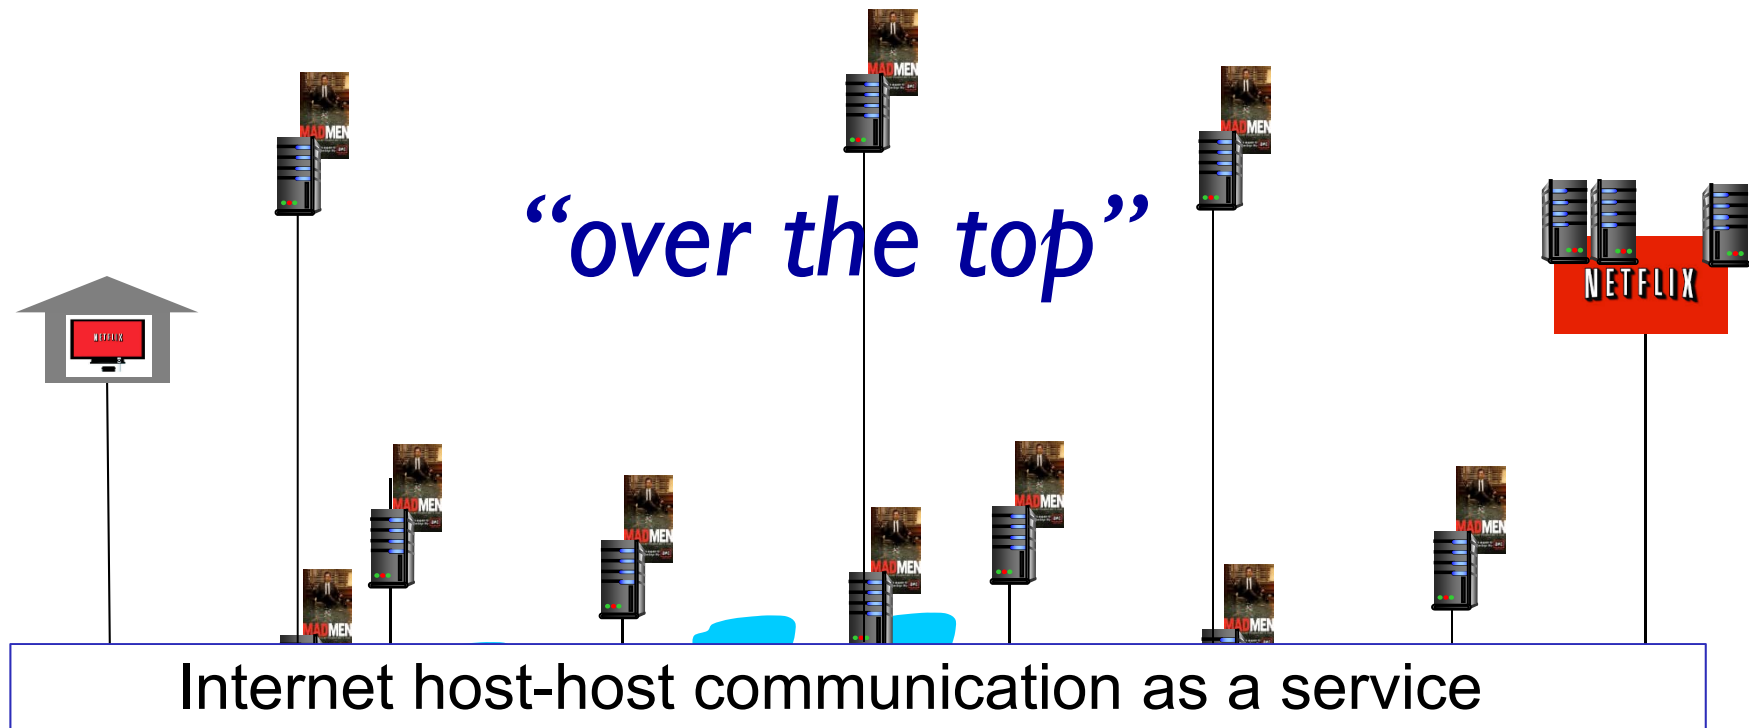
- *challenge*: how to stream content (selected from millions of videos) to hundreds of thousands of simultaneous users?
- *option 2*: store/serve multiple copies of videos at multiple geographically distributed sites (*CDN*)
  - *enter deep*: push CDN servers deep into many access networks
    - close to users
    - used by Akamai, 1700 locations
  - *bring home*: smaller number (10's) of larger clusters in POPs near (but not within) access networks
    - used by Limelight

# Content Distribution Networks (CDNs)

- CDN: stores copies of content at CDN nodes
  - e.g. Netflix stores copies of MadMen
- subscriber requests content from CDN
  - directed to nearby copy, retrieves content
  - may choose different copy if network path congested



# Content Distribution Networks (CDNs)



**OTT challenges:** coping with a congested Internet


- from which CDN node to retrieve content?
- viewer behavior in presence of congestion?
- what content to place in which CDN node?

*more .. in chapter 7*

# CDN content access: a closer look

Bob (client) requests video <http://netcinema.com/6Y7B>

- video stored in CDN at <http://KingCDN.com/NetC6y&B23V>



# Case study: Netflix

